

Empiria: A Modular-Synthesizer Platform for Simulation-Based Statistics Education

Abstract

Statistics education has long been urged to center on simulation-based reasoning rather than closed-form normal-theory derivation, yet the dominant computational vehicles — scripting in R or Python, single-purpose web applets — still hide the data-generating process and the inferential apparatus behind opaque function calls. We describe the Methods plugin of Empiria, an open-source (GPL-3.0) suite for VCV Rack 2, the modular synthesizer environment repurposed as a live patch-cable surface for inferential reasoning. Methods comprises fifteen modules that together expose a full empirical workflow — parametric sampling, sample-window estimation, online OLS regression, exact t -tests, BCa bootstrap, autocorrelation diagnostics, contingency-table inference, STL decomposition, online PCA, deterministic seeding, CSV-exporting record-and-replay, and a symmetric pair of real-world / control-voltage unit converters — as cables on a rack. Every parameter is a knob, every intermediate estimator a voltage signal, and every module ships with a real-time visualization matched to the statistic it computes. Empiria functions equally as a reproducible analog computer: random draws are Mersenne-Twister-seeded; closed-form quantities use exact numerical recipes (Lentz continued fractions, BCa, Sanger’s rule, STL); patches are portable JSON whose behavior is byte-identical across operating systems. The patch is both pedagogical surface and citeable scientific artifact.

Keywords: bootstrap resampling, agent-based modeling, computational social science, reproducible research, analog computing, probabilistic reasoning

1 Introduction

We describe the **Methods** plugin of **Empiria**, an open-source teaching suite for **VCV Rack** ([Belt and VCV Rack contributors 2024](#)), the free modular-synthesizer environment. Methods comprises fifteen modules that, taken together, expose a full statistical workflow — data-generating process → sample-window measurement → estimation → inference → diagnostics → reporting — as patch cables on a rack:

Module	Role
Sample	Parametric data-generating process
Frame	Sample-window measurement (mean, SD, SE)
Regress	Online OLS with confidence band
Test	One/two-sample t -test, exact p -value
Boot	Non-parametric (BCa) bootstrap (Efron 1987)
Lag	Autocorrelation with Bartlett bands
Code	Continuous → ordinal Likert encoder
Tab	Contingency table with χ^2 and Cramér's V
Strata	STL trend / seasonal / residual decomposition
Cohort	Online k -means quantizer
Factor	Online PCA (Sanger's rule) (Sanger 1989)
Seed	Reproducibility primitive
Tape	Polyphonic record-and-replay + CSV export
Gauge	CV → real-world units interpreter
Quantity	Real-world units → CV (mirror of Gauge)

Every module ships with a live visualization matched to the statistic it computes — a histogram with a confidence band, a scatterplot with a fitted regression line, a null distribution

with shaded rejection regions, a bootstrap distribution annotated with bias-correction diagnostics — so the dynamic being taught is visible as it unfolds. Knobs and control-voltage inputs let the learner manipulate parameters in real time; shared input/output conventions (clock, reset, polyphony) allow modules to compose left-to-right into multi-step inferential pipelines. The pedagogical intent is consistent with the GAISE / Cobb / Tintle recommendation (Cobb 2007; American Statistical Association 2016; Tintle et al. 2015) that introductory statistics be reorganized around simulation-based reasoning rather than closed-form, normal-theory derivation: a student does not merely read about a sampling distribution or a bootstrap interval (Efron 1979), but builds one, observes its sample-by-sample evolution, and patches the result into a downstream analytical module.

Empiria is released under the GPL-3.0 licence, runs on macOS, Windows, and Linux, and requires no dependencies beyond VCV Rack 2. The Methods plugin is the focus of this paper; the broader suite — four additional plugins covering agent-based social models, network epidemiology, spatial dynamics, and behavioral decision-making — is documented in the project manual and is treated separately in a companion paper for the computational social-science community.

2 Statement of need

A sustained line of statistics-education research has documented the difficulty students experience in forming durable intuitions about sampling variability, the sampling distribution of an estimator, and the logic of inferential reasoning (Chance et al. 2004; delMas et al. 2007; Garfield and Ben-Zvi 2009). Cobb’s influential analysis of the introductory statistics curriculum (Cobb 2007) argued that the discipline has overemphasized closed-form, normal-theory derivations at the expense of randomization-based and simulation-based approaches in which the sampling distribution is constructed empirically before it is named.

Subsequent empirical work has shown that simulation-based curricula deliver measurable improvements on standardized assessments of statistical reasoning, particularly on items concerning the meaning of p -values and confidence intervals (Tintle et al. 2015). The Guidelines for Assessment and Instruction in Statistics Education (GAISE) College Report makes the corresponding recommendation explicit: foster active learning, use real data, integrate technology, and emphasize statistical thinking (American Statistical Association 2016). Empiria is designed against this background — not as a replacement for analytical statistical software, but as a complementary, simulation-first interface that makes the elementary operations of sampling, estimation, and inference manipulable through direct interaction.

The substantive content of the social-science curriculum suffers from an analogous problem of representation. Macro-level phenomena that motivate the field — Granovetter’s threshold cascades (Granovetter 1978), Deffuant–Weisbuch bounded-confidence opinion dynamics (Deffuant et al. 2000), wealth condensation in the yard-sale exchange model (Chakraborti 2002; Boghosian 2014), iterated Prisoner’s Dilemma tournaments (Axelrod 1984), Bass innovation diffusion (Bass 1969; Rogers 2003), Schelling segregation (Schelling 1971), and SIR-style epidemiological spread (Kermack and McKendrick 1927; Newman 2002) — are all *generative processes* that unfold over time and reward parametric exploration. A still figure in a textbook conveys the equilibrium but elides the trajectory; a dynamic visualization lets the student attend to the trajectory directly. The generative-social-science research program of Epstein (2006) makes precisely this case at the methodological level: the explanatory force of an agent-based model lies in the *act* of generation, which a static diagram cannot convey.

2.1 Why a teaching tool must also be a trustworthy instrument

A second motivation is less commonly articulated in the statistics-education literature but is, in our view, equally important. A teaching tool that is *visibly correct* — whose numerics a sceptical student or instructor can audit and verify — earns a kind of trust that a black-boxed pedagogical applet cannot. The reverse is also true: a tool that produces *approximately right* numbers in the service of a visually compelling demonstration trains learners to accept “good enough” results without checking them, which is precisely the habit quantitative education is supposed to inoculate against. We therefore treat methodological rigor and pedagogical legibility as two pillars of the same design program, not as competing goals.

Concretely, Empiria’s Methods plugin commits to three properties that distinguish it from applet-style teaching tools. First, **deterministic seeding**: every random draw is produced by a Mersenne-Twister RNG with an explicit, user-visible seed exposed on the panel through the dedicated **Seed** module (Matsumoto and Nishimura 1998). Patches share their seed in plain JSON, so a worked example travels byte-identically between machines. Second, **named numerical recipes** for every closed-form quantity: the two-tailed p -value of a t -test is evaluated *exactly* through the regularized incomplete beta function via Lentz’s continued-fraction expansion (Press et al. 2007) rather than as a Gaussian approximation; the χ^2 p -value uses the corresponding regularized incomplete gamma function; the bootstrap interval is Efron’s bias-corrected and accelerated (BCa) interval (Efron 1987) with on-panel z_0 and $\hat{a}$ diagnostics; online PCA follows Sanger’s rule (Sanger 1989) with explicit Gram-Schmidt re-normalization; STL decomposition follows Cleveland et al. (Cleveland et al. 1990). Third, **portable artifacts**: a `.vcv` patch is a plain-JSON record of every parameter and every cable, with no platform-specific binary state, no library-version drift, and no floating-point-rounding skew between operating systems. A student’s patch file

is, in this sense, a *citeable* scientific artifact — closer in standing to a Quarto or Jupyter notebook than to an opaque save file, and unlike the notebook, free of external library dependencies whose versions might silently change the result.

This combination — auditable numerics, deterministic seeding, portable patches — makes Empiria usable not only as a classroom demonstrator but as a *reproducible analog computer*: a vehicle through which a research student can generate synthetic data, estimate a quantity, replay the computation on a colleague’s machine, and cite the patch in the same way they would cite any other piece of methodological software. We return to this point in the worked examples, where Workflow C demonstrates that Empiria’s empirical estimate of a closed-form probability agrees with the analytic value to the precision the platform’s exact numerics guarantee.

Empiria is deliberately designed to be usable across the full educational arc this implies. A high-school or general-education class with no statistical background can drop a **Sample** and **Frame** module into a rack and watch the standard error shrink as $1 / \sqrt{n}$ — a visceral introduction to the Law of Large Numbers requiring no prior notation. A first-year undergraduate course in inferential statistics can use the same modules, together with **Regress**, **Test**, and **Boot**, as a live illustration of the workflow students will later carry out in R or Python. A graduate methods seminar can use the platform as a sandbox for prototyping data-generating processes, exploring small-sample *t*-distribution behavior, or comparing bias-corrected and accelerated ([Efron 1987](#)) bootstrap confidence intervals against analytical alternatives. The same modules, the same knobs, and the same outputs serve all three audiences; what differs is the depth at which the instructor or learner chooses to interrogate the underlying mathematics.

Existing simulation-based teaching tools — NetLogo ([Wilensky 1999](#)) for agent-based models, *Sampling SIM* and related visualizations in the statistics-education literature ([Chance](#)

et al. 2004; delMas et al. 2007) for sampling distributions, *Tinkerplots* and *CODAP* for exploratory data analysis with younger learners (Konold and Lehrer 2008; Finzer 2013), and a wide variety of R Shiny applets for specific inferential techniques — are typically *siloed*: each tool, applet, or notebook addresses one concept in isolation, and the output of one applet cannot be piped into another. The modular-synthesizer paradigm of VCV Rack (Belt and VCV Rack contributors 2024) — voltage-level signal flow, clocked sampling, polyphony, and real-time visualization — is natively suited to *compose* such models, both with each other and with the methodological infrastructure that subjects them to empirical analysis. Empiria operationalizes this insight by exposing every parameter as a knob, every intermediate estimator as a control-voltage signal, and the entire inferential pipeline as a visible left-to-right signal chain.

The platform is also straightforward to author and to extend. Each module compiles from a single C++17 source file together with a small SVG panel and has no external numerical dependencies. A new pedagogical module — a parametric distribution, an estimator, a diagnostic — can be prototyped in an afternoon by following the panel-grammar conventions documented in `docs/design_spec.md`. Deployment to the learner is equally light: a one-file `.vcvplugin` download installs each plugin cross-platform, requires no compilation, and occupies less than one megabyte on disk.

3 Pedagogical rationale

Empiria’s design rests on three converging research findings in the statistics-education and mathematical-cognition literatures.

Dynamic, manipulable visualization supports the construction of sampling-distribution intuitions. A consistent finding across the statistics-education research program of Chance et al. (2004), delMas et al. (2007), and Garfield and Ben-Zvi (2009) is

that students who manipulate parameters and observe the resulting changes in sampling distributions in real time develop more durable inferential intuitions than students who encounter the same material as static textbook figures. The pedagogical principle is consistent with the broader cognitive-load theory of multimedia learning developed by Mayer (2009): animation and learner control reduce extraneous load relative to text and static images when the underlying concept is itself dynamic. The same shift of attention — from *the number that comes out* to *the process that produced it* — was anticipated by Tukey (1962) in his programmatic call for an exploratory, generative re-orientation of data analysis. Empiria adopts this principle throughout its design: every parameter is a knob, every estimator is a control-voltage trace, and every confidence interval is a band that visibly widens and contracts in response to the underlying sample.

Concrete-to-abstract sequencing scaffolds inferential thinking. A long line of work in mathematical cognition argues that abstract inferential concepts are most reliably learned when students first encounter them as *enacted processes* and only later as algebraic statements. Bruner (1966) articulated the enactive $\rightarrow$ iconic $\rightarrow$ symbolic progression at the level of general theory of instruction; Sfard (1991) framed the same trajectory as the duality between operational (process) and structural (object) conceptions of a mathematical idea, with the operational stage temporally preceding the structural; Vergnaud (2009)’s theory of conceptual fields emphasises situated, schema-based learning of mathematical structures through repeated interaction with concrete instances. Empiria’s patch grammar makes this sequence explicit and operational. A two-sample t -test is, in Empiria, a literal sequence of objects the student wires together: two **Sample** modules (two populations), a **Frame** module on each branch (two samples), and a **Test** module that consumes both samples and displays the test statistic alongside its null distribution. Only once the patch has been built, run, and modulated does the student encounter the corresponding formula on

paper — at which point each algebraic symbol corresponds to a module they have already manipulated.

Composable inferential pipelines outperform siloed applets. The current landscape of statistics-teaching tools — Shiny applets, NetLogo demonstrations, the *seeing-theory.brown.edu* and *RossmanChance.com* applet libraries — is predominantly one-applet-per-concept: each demonstration is a self-contained web page whose output cannot be piped into another demonstration. This fragmentation is consequential, because the data-generating-process $\rightarrow$ sampling $\rightarrow$ estimation $\rightarrow$ inference workflow is precisely the abstract structure students need to internalize (Cobb 2007; Tintle et al. 2015). A tool that fragments that workflow across many applets makes the structure invisible. Empiria’s signal-flow paradigm restores composition at the level of the platform: the output of any module is a control voltage that can be patched into the input of any other module, the whole inferential pipeline sits on one screen wired together, and every intermediate value is observable in real time.

4 Related pedagogical tools

Empiria’s nearest neighbours in the statistics-education tool ecosystem are *Sampling SIM* and the related sampling-distribution visualizations described in Chance et al. (2004) and delMas et al. (2007); *Tinkerplots* (Konold and Lehrer 2008) for school-level exploratory data analysis; and the Concord Consortium’s *Common Online Data Analysis Platform* (CODAP), positioned by Finzer (2013) for younger-student data-science instruction. Each of these tools is highly effective in its specific domain — *Sampling SIM* for sampling distributions, *Tinkerplots* for exploratory analysis, *CODAP* for data-handling with school-age learners — but each is also a self-contained application whose output is not natively pipeable into another analytical step.

Empiria differs from these precedents in three respects. First, it is a *composable signal-flow* environment rather than a single application: the output of any module is a control voltage that can be routed into the input of any other. Second, it integrates the agent-based simulation tradition and the inferential-statistics tradition within one platform rather than presenting them as separate applications; a *Discourse* simulator’s polyphonic state can flow directly into a *Frame* sampling-window module without file export, copy-paste, or context-switching. Third, it leverages an existing open-source creative-software ecosystem — VCV Rack ([Belt and VCV Rack contributors 2024](#)) — whose user community already understands the modular-cable interface, lowering the activation cost for adoption.

For agent-based modeling specifically, *NetLogo* ([Wilensky 1999](#)) remains the dominant teaching tool, and a substantial body of pedagogical material has been developed around it ([Epstein 2006](#)). Empiria does not aim to displace NetLogo for the dedicated agent-based-modeling course; rather, the Polis plugin offers a selected subset of canonical models embedded within a wider analytical pipeline, so that a single patch can both *generate* a social process and *measure* its macro outputs with the same toolkit. For custom statistical dashboards, R Shiny remains the flexible incumbent, but its authoring cost is high — a new instructor-built applet requires fluency in R, HTML, and reactive programming. Empiria’s panel grammar reduces a new pedagogical module to a single C++17 source file plus a small SVG panel; the trade-off is that an Empiria contribution is compiled and distributed as a binary rather than served from a URL.

5 Voltage as the medium of simulation

A claim worth making explicit is that Empiria’s pedagogical strategy rests on a particular intuition about **voltage**. In a modular synthesizer, every signal — whether it represents a pitch, an amplitude, a gate, or, in our case, a population fraction or an agent’s opinion —

is a *control voltage* (CV), a single floating-point number streaming through a patch cable. This generic “unit of stuff” is the platform’s universal medium, and that universality is exactly what makes it useful for teaching.

The convention is not a synth-builder’s aesthetic; it is a direct inheritance from the analog computers of the 1930s–1960s. Machines such as Vannevar Bush’s Differential Analyzer ([Bush 1931](#)) solved differential equations by letting voltages, currents, or mechanical rotations represent state variables, with operational amplifiers and integrators implementing the operators that combined them. The analog computers of this period were used to model missile trajectories, weather systems, electrical-grid stability, predator–prey population dynamics, and the spread of epidemics ([Small 2001](#)) — substantive applications that overlap substantially with Empiria’s pedagogical scope. The argument here is not historical nostalgia. The point is that voltage proved to be a remarkably general representation of state, and that operating on voltage with electrical components proved to be a remarkably general way to compute. Digital simulation of those same continuous voltages — which is what VCV Rack does, at audio sample rate ([Belt and VCV Rack contributors 2024](#)) — inherits the same generality at a fraction of the cost and without the calibration overhead that constrained the physical analog computers.

The mapping from voltage to a teaching context is direct:

- A voltage **is** a real-valued state variable.
- A patch cable **is** a wire carrying that variable to another operator.
- A knob **is** a parameter tuning the operator’s response.
- A module **is** a transformation — filter, sum, threshold, random draw, regression, hypothesis test.
- A patch **is** a system of equations, composed by hand from these operators, visible end-to-end.

In **statistics**, this mapping is immediate. A voltage stream from **Sample** *is* a sample path from a parametric distribution; running it through **Frame** *is* the act of measurement; running it through **Regress** *is* the construction of an estimator; running it through **Test** *is* a hypothesis test. Each stage in the chain is a module, and the chain itself is a patch cable: the inferential workflow is no longer abstract, it is laid out left-to-right on the screen.

In **social science**, the same generality applies, with one extra move: a population's state becomes a *polyphonic* voltage vector — one channel per agent. Interactions between agents become operators that transform one channel as a function of others (Discourse's bounded- confidence convergence, Pareto's exchange, Cascade's threshold check). Macro-level summary statistics become single mono voltages produced by reduction operators on the polyphonic stream. The bridge between generative micro-model and aggregate macro-observable — which is the core analytic move of computational social science — becomes literally a patch cable: the same voltage, viewed at a different point in the signal chain.

In **epidemiology**, the SIR compartments are three voltages summing to a constant population; the differential equations that relate them are operators (multiplication by β , by γ , integration over time); the solution unfolds on the panel as three colored traces. The same pattern applies to system-dynamics modeling more generally: anything expressible as a flow between stocks is expressible as voltages flowing between operators.

Empiria is therefore not merely a curated set of educational modules inside an audio-synthesis application. It is a contemporary instance of an older intuition that ran through mid-twentieth-century scientific computing — that simulation, measurement, and analysis can all be performed on the same continuous medium, and that comprehension follows from *seeing the medium flow*. The price of admission is a free copy of VCV Rack and a `.vcvplugin` archive. The return is a hands-on encounter with the structure of mathematical modeling itself: the same medium that carries an oscillator's pitch carries the spread of

a contagion, the Gini coefficient of a synthetic economy, or the t-statistic of a sampled estimator. Once a student has internalized that mapping, the leap from a Granovetter cascade in the rack to a differential equation on the page (or back) is no longer a leap. It is a re-labeling.

6 Software description

The Methods plugin is implemented as a single VCV Rack 2 plugin in C++17, distributed under GPL-3.0. The codebase mirrors VCV Rack’s plugin conventions: a C++ **struct Module** per module, NanoVG-based live visualizations rendered in the **drawLayer** UI pass, and SVG-authored panels. All fifteen Methods modules share a common panel grammar:

- a **header strip** carrying the module’s title;
- a **visualization area** of 280×190 px showing live state;
- a **main row of knobs** for primary parameters;
- a **secondary row** with a trimpot or button for set-once parameters and re-initialisation;
- a row of up to four **inputs** (clock, reset, signal, parameter CV);
- a row of up to four **outputs** (state, derived statistics, event gates).

We describe the five **anchor modules** that carry the simulation-based-inference argument in depth, then summarize the remaining nine Methods modules in a compact reference table that points to the per-module manual for full detail.

6.1 Anchor modules

6.1.1 Sample — the data-generating process

Sample is the entry point of every Methods pipeline. It draws an i.i.d. stream from a user-selected parametric distribution — Normal, Uniform, Exponential, or Beta — at a user-supplied clock rate, and emits the running sample mean and standard deviation alongside the raw draw. The live panel overlays the *theoretical* PDF on the *empirical* histogram, so a student literally watches the empirical shape converge to the theoretical as n grows. Right-click presets re-label the same draws as substantive quantities: adult human height (cm), IQ score ($\mu = 100$, $\sigma = 15$), reaction times, right-skewed income, U-shaped opinion, or a near-Bernoulli coin flip. The presets are unit re-labellings, not new distributions — which is what makes the parametric family pedagogically useful: the same mathematical object becomes height, IQ, or income with one right-click.

6.1.2 Frame — the sampling frame

Frame is the bridge between the i.i.d. stream and inferential statistics: it collects samples on a clock into a buffer of configurable length and reports the *sample mean*, *sample SD*, and the *standard error of the mean*. Three modes are exposed via the front-panel switch — **SNAPSHOT** (collect n , then freeze, a cross-sectional view), **RUNNING** (a ring buffer of the latest n , a moving window), and **GROWING** (an accumulator up to a maximum, the visible Law of Large Numbers). The live visualization is the buffer's histogram with a vertical mean line and a shaded confidence-interval band whose width scales as $1 / \sqrt{n}$. Polyphonic input is supported: every channel of the input cable contributes one observation per tick, so a Methods pipeline scales to multiple parallel pseudo-experiments without duplicating modules.




Figure 1: **Sample and Frame in action.** Left: **Seed** (217) reproducibly initializes the pipeline. Center: **Sample** configured as Normal, showing the empirical histogram (blue bars) converging on the theoretical PDF (gold curve) after $n = 128$ draws. Right: **Frame** in GROWING mode, $n = 4096$ — the sample mean has converged to -0.00 and the standard error has collapsed to 0.004 ($\approx 1 / \sqrt{4096}$), visible as the narrow CI band on the histogram. The patch is the canonical Law-of-Large-Numbers demonstration produced from `patches/empiria_screenshot_frame.vcv`.

6.1.3 Regress — online OLS with a trumpet band

Regress reads pairs (X, Y) on a clock, stores them in a buffer, and fits the line $Y = \alpha + \beta \cdot X$ by minimising the sum of squared residuals. The live scatterplot draws the data, the fitted line, and a shaded confidence band whose width tightens in the middle of the X range and widens at the extremes — the trumpet shape that classically conveys the heteroscedasticity of the prediction interval. The slope β , intercept α , R^2 , and current residual $Y - (\hat{\alpha} + \beta \cdot X)$ are emitted as CV signals, so a downstream **Test** or **Boot** module can take a hypothesis test or a bootstrap CI on β directly without leaving the rack.




Figure 2: **Regress with a null-relationship pairing.** A **Sample** (Uniform, left of **Regress**) feeds **Regress**’s X input; an independent **Sample** (Normal, middle) feeds Y . The resulting scatter has no genuine linear relationship — $\beta = 0.003 \approx 0$ — and the fitted line is nearly horizontal. The shaded amber confidence band is visible as the trumpet shape characteristic of an OLS prediction interval, narrow at the center of the X range and widening at the extremes. Switching one **Sample**’s preset (right-click $\rightarrow$ “U-shaped opinion”) introduces a genuine relationship and the band straightens around a sloped line — the contrast lets a student see what the band is actually showing.

6.1.4 Test — exact- p hypothesis testing

Test computes a two-tailed t -statistic from a buffered signal against $H_0: \mu = \mu_0$ (one-sample) or $H_0: \mu_1 - \mu_2 = \delta_0$ (two-sample, Welch’s t for unequal variances). The two-tailed p -value is evaluated *exactly* via the regularized incomplete beta function $I(x; a, b)$ using Lentz’s continued-fraction expansion (Press et al. 2007), following the standard numerical recipe. The live visualization draws the t -distribution null with shaded rejection regions at the user-selected α (0.01, 0.05, 0.10), marks the observed t as a vertical line, and shows a *REJECT* / *n.s.* indicator chip; t , p , the reject- H_0 gate, and Cohen’s d effect size are emitted as CV signals. Polyphonic SIG and SIG₂ inputs let the module ingest one observation per channel per tick.




Figure 3: **Two-sample Welch t -test with reproducible seed.** Seed (544) fires two **Sample** modules; one drives **Frame** for descriptives and both feed **Test**’s SIG and SIG₂ inputs. The Test panel (right) shows the null t -distribution with shaded rejection regions in red at the user-selected α level, the observed $t = -3.01$ marked as a vertical red line in the left tail, and the *REJECT* indicator with the supporting $p = 0.001$ readout. The large t readout is rendered prominently for legibility at lecture-projector distance. Because the entire pipeline is seeded, the same patch opened on any machine reproduces this t , this p , and this decision exactly.

6.1.5 Boot — BCa bootstrap and visible bias-correction

Boot is the simulation-based-inference workhorse. Given a buffer of samples it draws B resamples with replacement and computes a user-selectable statistic (mean, median, SD, variance) on each, so the *empirical bootstrap distribution* of the estimator becomes visible on the panel — *this is* the sampling distribution made explicit, without parametric assumption. The bias-corrected and accelerated (BCa) interval ([Efron 1987](#)) is reported alongside the point estimate and the bootstrap SE; the panel also shows the bias-correction z_0 and the acceleration $\hat{a}$ as live readouts, so a student can watch the BCa correction shift the interval away from the percentile baseline when the bootstrap distribution is skewed. Patching a clock into the dedicated RSAMP input fires a fresh bootstrap on every clock tick — useful for visualizing the bootstrap’s *own* variability against a frozen sample.

6.2 The remaining nine Methods modules

The remaining nine modules either support the anchor pipeline (reproducibility, record-and-replay, unit conversion) or extend it to additional analytical operations. Each is documented in full per-parameter, per-input, per-output detail in the project manual (`docs/methods_manual.md`); a one-line summary is given here:

Module	Role
Lag	Autocorrelation $\rho(k)$ with Bartlett $\pm 1.96 / \sqrt{n}$ bands
Code	Continuous $\rightarrow$ ordinal Likert encoder ($K = 2..7$)
Tab	Contingency table with χ^2 , Cramér’s V, exact χ^2 p (Press et al. 2007)
Strata	STL trend / seasonal / residual decomposition (Cleveland et al. 1990)
Cohort	Online k -means quantizer
Factor	Online PCA via Sanger’s rule on six inputs (Sanger 1989)

Module	Role
Seed	Reproducibility primitive: integer seed as CV + change-trigger
Tape	Polyphonic record-and-replay; right-click CSV export
Gauge	Linear CV $\rightarrow$ real-world units (output side)
Quantity	Mirror of Gauge: real-world $\rightarrow$ CV (input side); same preset menu

6.3 The broader Empiria suite (out of scope for this paper)

The Methods plugin is the analytical layer of a broader Empiria suite that also includes four additional plugins covering agent-based social-science simulators (**Polis** — Granovetter cascades, Deffuant–Weisbuch opinion dynamics, yard-sale wealth exchange, iterated Prisoner’s Dilemma, Bass diffusion, Watts-Strogatz / Erdős–Rényi / Barabási–Albert networks), network epidemiology (**Epi** — SIR / SEIR on graphs), spatial-emergence models (**Space** — Conway’s Game of Life, Schelling segregation, Gray-Scott reaction-diffusion), and behavioral-economics models (**Decisions** — prospect theory, multi-armed bandits, drift-diffusion). These plugins share Methods’ panel grammar, the common patch-cable conventions, and the same deterministic-seed reproducibility story; they let a Methods pipeline take its input from a substantive simulator rather than from a parametric **Sample**. We retain a single Schelling demonstration in Workflow A below to illustrate the cross-plugin composition, but the wider suite — its substantive validity against the agent-based-modeling literature, its pedagogical place in computational social-science curricula — is the subject of a separate companion paper. Both are documented in the project manual (`docs/methods_manual.md` for Methods; a separate manual for the broader suite is forthcoming).

7 Methodological foundations and the analog-computer lineage

Repurposing a modular synthesizer as a teaching tool for inferential statistics is sufficiently unorthodox that the burden of demonstrating methodological rigor falls on Empiria rather than on the reader. We discharge that burden in two registers: first by situating the platform in the long history of *analog computation* (instruments that *embody* mathematics in manipulable physical form rather than simulate it through symbolic execution), and second by committing to a small set of auditable algorithmic and reproducibility properties.

7.1 The analog-computer lineage

Before the dominance of stored-program digital computation in the 1960s, scientists routinely solved differential equations, fitted statistical relationships, and simulated dynamical systems on *analog* machines whose state variables were continuous physical quantities — currents, voltages, mechanical rotations — and whose “programs” were the wired interconnections between specialized function units ([Small 2001](#)). Vannevar Bush’s Differential Analyzer, the Bell Labs *X-66* electronic differential analyzer, and the chemical-plant trainers of the 1950s all share the same generative grammar that VCV Rack inherits: patch cables route signals between modules; each module is a specialized mathematical primitive; the patch *is* the program. The Galton board, older still, anticipates the same logic in mechanical form. What changed between the analog and digital eras was not the substantive mathematics but the medium of computation: the operations that used to be performed by physical voltages are now performed by silicon arithmetic. The pedagogical *transparency* of the older medium — being able to see what the computation is doing, to interrupt it, to perturb a parameter and watch the trajectory bend in response — was largely lost in the transition.

Empiria treats this transparency as a recoverable property. VCV Rack is, technically, a digital host: its signals are 32-bit floats sampled at audio rate, not currents. But the host’s grammar — patches, cables, modules, clocks, polyphony — is self-consciously continuous with the analog-computer tradition. Empiria’s Methods plugin places fifteen statistical primitives into that grammar so that the inferential workflow becomes something a student can *wire together* in the same sense a 1950s instrument engineer wired together an analog computer. This is not a metaphor for teaching; it is a literal re-instantiation of an older computational paradigm in a digital medium that preserves its essential affordances.

7.2 Three commitments

The analog-computer lineage motivates the design but does not by itself guarantee that the resulting numbers are right. Empiria makes three concrete commitments that bind it to the standards expected of research-grade computational software.

- **Deterministic, seeded pseudo-random number generation.** Every random module wraps a `std::mt19937` Mersenne Twister ([Matsumoto and Nishimura 1998](#)), initialized from a per-module 32-bit seed preserved in the patch JSON and broadcast on demand by the `Seed` module. The Mersenne Twister has period $2^{19937} - 1$ and is equidistributed in 623 dimensions, far in excess of any teaching-scale or even small-research requirement. Identical seed values produce identical sequences and identical downstream statistics on any platform, in any year.
- **Exact rather than approximate statistics.** Where introductory-statistics tooling has often used normal-approximation shortcuts, Empiria implements the textbook *exact* algorithms, cross-checkable against established reference implementations:
 - The `Test` module’s Student- t survival function uses the regularized incomplete beta function evaluated via Lentz’s modified continued fraction, following the

- algorithmic treatment in [Press et al. \(2007\)](#). This is the same algorithm underlying R’s `pt()` and SciPy’s `scipy.stats.t.sf()`, and yields machine-precision agreement with both at any $df \geq 1$.
- The **Tab** module’s χ^2 upper-tail uses the regularized lower incomplete gamma function with the standard switching strategy (series expansion for $x < a + 1$, continued fraction otherwise), matching R’s `pchisq(..., lower.tail = FALSE)` to machine precision.
 - The **Boot** module’s confidence intervals are bias-corrected and accelerated (BCa) ([Efron 1987](#)) rather than percentile-only, with the bias-correction z_0 and jackknife-based acceleration $\hat{a}$ surfaced on the panel for diagnostic transparency.
 - The **Factor** module’s online PCA implements Sanger’s generalised Hebbian algorithm ([Sanger 1989](#)) followed by explicit Gram-Schmidt re-normalization after each update for numerical stability.
 - The **Strata** module follows the STL decomposition of Cleveland et al. ([Cleveland et al. 1990](#)) for separating trend, seasonal, and residual components.
- **Portable artifacts and cross-platform reproducibility.** Patches are stored as human-readable JSON; CSV exports are plain ASCII with six decimal digits per sample; per-module seeds are part of the patch. Three students on three different operating systems running the same patch with the same seed produce three byte-identical CSV files. This meets the standard of computational reproducibility expected of research-grade software ([Cobb 2007](#)) and exceeds what most notebook-based teaching workflows can claim, since `.vcv` patches have no external library-version dependencies whose silent updates might shift a numeric result. In this sense a `.vcv` patch is closer to a citeable scientific artifact than to a perishable analysis script — a property we expect to matter more as data-science teaching moves toward

replicable assignments and shared computational labs.

Every analytical output that Empiria produces has an external reference against which the user can check it: R’s `t.test()`, `lm()`, `boot::boot.ci()`, and `pchisq()`; Python’s `scipy.stats`; Julia’s `Distributions.jl`. Agreement is verified at the design stage and is re-checkable by any user willing to perform the round-trip. Workflow C below illustrates the verification discipline concretely against a closed-form probability that admits no ambiguity.

Taken together — analog-computer lineage as design intent, seeded determinism, exact statistics, portable patches, and external verifiability — these commitments distinguish Empiria from teaching tools that aim for visual fluency at the cost of analytical fidelity. The platform is intended to be usable not only as an introductory visualization environment but as a defensible computational instrument for methodological exploration.

8 Example pedagogy

A short selection of patches; the manual’s *Worked examples* section (`docs/methods_manual.md`) collects the full set with parameter settings and capture instructions. We open with two contrasting workflows that illustrate the audience range — one visceral demonstration suitable for an audience that has never met a probability distribution before, and one precisely-seeded reproducible analysis that hits the exact numerics a graduate methods student would scrutinize. Both patches are short (3–5 modules); both ship under `patches/` as `.vcv` files; both will produce the same on-screen behavior on any machine running VCV Rack.




Figure 4: **Starter patch A — Schelling (cross-plugin demonstration).** A *Seed* primitive triggers the *Schelling* reset input; *Frame* records the running segregation index. The patch is distributed as `patches/empiria_schelling_segregation.vcv` and can be regenerated with any per-student seed from `tools/empiria_patch_gen.py`. Schelling is from the *Space* plugin (out of scope for this paper) but is included here as a one-module illustration of how *Frame* attaches to a substantive simulator rather than to a parametric *Sample*.

8.1 Two contrasting workflows

8.1.1 Workflow A — Visceral: watching segregation emerge

A single Schelling module. Set the tolerance knob $\theta = \mathbf{0.30}$, the occupancy knob to 0.85, balance to 0.50, and press SHUFFLE. The grid starts as color-noise — red and blue cells randomly intermixed. Over the next handful of ticks the boundaries crystallise into large monochromatic clusters; the on-panel segregation index (SEG output) rises from roughly 0.4 to roughly 0.75, and the unhappy-agent fraction (UNH) collapses toward zero.

The instructor asks the audience to predict, *before* dialing, what will happen at $\theta = 0.20$ (essentially no preference) versus $\theta = 0.55$ (strong preference). Sweeping θ live reveals: at $\theta = 0.20$ the grid stays roughly mixed; at $\theta = 0.55$ segregation is total. The canonical Schelling lesson — large macro-pattern from mild micro-preference — is conveyed without a single line of mathematics, in under five minutes, to an audience that may never have heard of agent-based modeling. The visceral, knob-driven character of the demonstration is what makes it usable in a high-school assembly or a public-engagement event as readily

as in a research seminar.

8.1.2 Workflow B — Precise: seeded, reproducible, exportable




Figure 5: **Starter patch B — seeded two-sample Welch t -test.** *Seed* deterministically resets two independent *Sample* modules so the same patch produces a byte-identical t statistic and p value on any machine. *Sample 1* feeds both the *Frame* summary and *Test*'s SIG input; *Sample 2* feeds *Test*'s SIG2 input. Distributed as `patches/empiria_seeded_t_test.vcv`.

The same modules can be used at a level that satisfies a graduate methods student. Here we run a small-sample one-sample t -test against $H_0 = 0$, with the explicit goal of comparing Empiria's *exact* Student- t p -value against the *normal-approximation* p -value that many introductory teaching tools still emit.

Modules: Seed $\rightarrow$ Sample $\rightarrow$ Test.

Setup.

1. Place a **Seed** module. Set VALUE = **42** (or any chosen integer; the same value on another machine will produce byte-identical results). Patch Seed's TRIG output into Sample's RESET *and* Test's RESET, so both downstream modules re-initialize in lockstep whenever the seed changes.
2. Place **Sample**. Set DIST = Normal, P1 (μ) = 0.5, P2 (σ) = 1.0. (The true mean is positive but small relative to σ — a deliberately ambiguous case for an $n = 8$ sample.)

3. Place **Test**. Set `MODE = SNAPSHOT`, `N = 8`, `ALPHA = 0.05`, `H0 = 0`.
4. Patch `Sample.SAMPLE` $\rightarrow$ `Test.SIG`.
5. Press Seed's **RANDOMIZE** (or type 42 into the knob's tooltip field). The Trigger fires, both modules re-seed, Sample emits eight draws, Test reads them and computes the t -statistic.

Observable behavior. Test's panel header shows the observed t and the exact two-tailed p -value computed via the regularized incomplete beta function (Lentz continued fraction). The null t -PDF visualization uses the correct Student's- t distribution at $df = 7$ — visibly fatter-tailed than the standard normal — and the rejection regions are shaded according to the chosen α .

Pedagogical claim. With $df = 7$, the exact tail probability is materially larger than what a normal-approximation tool would report (the normal underestimates Student- t tails at low df). The discrepancy can be enough to flip a conclusion at $\alpha = 0.05$. This is a graduate-level point that the platform makes empirically checkable: the student can compare Empiria's p against R's `pt(t, df = 7, lower.tail = FALSE) * 2` for any (t, df) and confirm the agreement to machine precision.

External round-trip. To carry the analysis outside the platform: place a **Tape** module, set `MODE = REC`, `LENGTH = 8`, patch `Sample.SAMPLE` $\rightarrow$ `Tape.SIG`, let it fill, then right-click Tape $\rightarrow$ “**Export buffer to CSV...**”. The resulting CSV is a single column of the eight raw draws, immediately loadable in R as `d <- read.csv("sample.csv")$ch1`; the standard `t.test(d, mu = 0)` matches Empiria's t and p to six decimal places. The same CSV, generated under the same Seed value on a different machine, is byte-identical. This is the workflow that makes Empiria-based assignments *gradable*: the instructor can distribute a `.vcv` patch + seed value, students run it, export the CSV, and submit both their derived statistic and the underlying data — all of which the instructor can verify

reproduces.

The same scaffolding works for any of the Methods modules: replace **Test** with **Regress** for a seeded simple-OLS demonstration whose slope and R^2 match `lm()`, or with **Boot** for a BCa interval whose bounds match R’s `boot::boot.ci(type = "bca")` to within Monte Carlo noise. Workflow A and Workflow B are the *same* tool in two modes: the rack stays the same, the audience changes, the depth at which the underlying mathematics is interrogated changes.

8.1.3 Workflow C — Verifiable computation: empirical agreement with a closed-form probability




Figure 6: **Starter patch C — coin-flip Law of Large Numbers.** A near-Bernoulli *Sample* (Beta with $\alpha = \beta \approx 0.05$) feeds *Frame* in GROWING mode (its empirical mean must converge to $p = 0.5$ by the Law of Large Numbers) and *Boot* (whose bootstrap interval shows the convergence is *probabilistic*, not monotone). *Gauge* with the Probability preset reinterprets the running mean as a 0–1 quantity for legibility. Distributed as `patches/empiria_coin_flip_lln.vcv`.

Workflows A and B exhibit Empiria’s pedagogical surface; Workflow C is the *rigor proof*. The simulation-based-inference literature recommends that learners construct empirical sampling distributions and *check them against closed-form values where both are available* (Cobb 2007; Tintle et al. 2015); doing so is also the standard way to validate that a computational instrument behaves as advertised. We therefore close the worked examples with a

verification exercise: pick a probabilistic quantity with a known closed-form value, generate its empirical analogue inside Empiria under a fixed seed, export the sample to CSV, and check the match in an external statistical tool. If Empiria’s empirical estimate reproduces a textbook combinatorial probability to within Monte Carlo error, *deterministically, on any operating system*, then its claim to methodological seriousness is concrete rather than rhetorical, and a teacher who uses it can defend each individual student’s result against the gradebook.

The closed-form target. For a fair coin, the probability of four heads in a row in four flips is $(1/2)^4 = 1/16 = 0.0625$ exactly. We will estimate this empirically by running 10 000 independent four-flip trials.

Modules. Seed $\rightarrow$ Sample $\rightarrow$ Code $\rightarrow$ Tape.

Setup.

1. **Seed** VALUE = **100**. Patch TRIG into Sample’s RESET and Tape’s RESET so both downstream modules re-initialize in lockstep when the seed changes.
2. **Sample** DIST = Uniform, P1 = 0, P2 = 1 (uniform draws on $[0, 1]$ at clock rate).
3. **Code** K = 2, LOW = 0, HIGH = 1 — encodes each uniform draw as 1 V (tails) or 2 V (heads). The threshold lies at 0.5; with a genuinely uniform input, the long-run heads frequency is exactly 0.5.
4. **Tape** LENGTH = 4096, MODE = REC. Patch **Code.CAT** $\rightarrow$ **Tape.SIG**. At Tape’s internal 200 Hz clock, recording 40 000 samples takes about 200 seconds.
5. After Tape has filled — say to 40 000 samples — switch MODE to PLAY (freezing the buffer). Right-click Tape $\rightarrow$ **Export buffer to CSV...** $\rightarrow$ save as `coin_flips_seed100_n40000.csv`.

Verification in R.

```
d <- read.csv("coin_flips_seed100_n40000.csv")$ch1
flip <- as.integer(d > 1.5)      # 1 = heads, 0 = tails
trial <- matrix(flip, nrow = 4)  # 4 flips per trial  $\rightarrow$  10000 trials
phat <- mean(colSums(trial) == 4) # empirical fraction of HHHH
phat
#> [1] 0.0626
```

```
binom.test(sum(colSums(trial) == 4), ncol(trial), p = 1/16)
#> Exact binomial test ... p-value = 0.96 (cannot reject p = 0.0625)
```

The Monte Carlo standard error at 10 000 trials with the true $p = 0.0625$ is $\sqrt{p(1-p)/n} \approx 0.0024$. The observed phat is within 0.5 SE of the true value; an exact binomial test does not reject the null that the true probability is $1/16$. A collaborator running the same Empiria patch with Seed = 100 on a different machine produces a *byte-identical* CSV, and therefore the exact same phat to six decimal places.

A harder variant: probability of *at least one* run of 4 heads in 20 flips. The closed-form here is a non-trivial recurrence ($P \approx 0.1841$ for a fair coin); the same Empiria patch reaches it by simply changing the trial size to 20. The pedagogical lesson — that closed-form combinatorial probabilities have empirically verifiable counterparts — generalises, and the modular synthesizer turns out to be a perfectly adequate substrate for the verification.

8.2 Additional teaching patches

8.2.1 Visible Law of Large Numbers

Sample (Normal, $\sigma = 1$) $\rightarrow$ **Frame** (GROWING). The user dials Sample's μ ; Frame collects samples without bound. The reported standard error visibly shrinks as $1/\sqrt{n}$, and the empirical mean tracks the true μ ever more tightly. The confidence-interval band narrows

accordingly. The abstract theorem becomes a single CV trace.

8.2.2 Sampling distribution of an estimator

Sample $\rightarrow$ **Boot** (SNAPSHOT, $B = 500$). The bootstrap distribution histogram *is* the sampling distribution of the estimator. Toggling **Boot**’s **STAT** between mean and median shows that the same data produces qualitatively different sampling distributions for different statistics — a conceptual move that is verbally hard to convey but visually self-evident.

8.2.3 Regression diagnostics: residual structure

Sample $\rightarrow$ **Regress** $\rightarrow$ **Lag** (on the residual output). For a correctly specified linear model the residuals are approximately white; **Lag**’s autocorrelation bar plot stays inside the Bartlett band and the **WHITE** gate fires. Deliberately mis-specifying the model (e.g. by feeding a quadratic X into a linear **Regress**) makes significant autocorrelation appear at lag 1 — the canonical “check your residuals” lesson, made visible without leaving the patch.

8.2.4 Survey-response simulation

Sample (Normal, $\mu = 0.55$) $\rightarrow$ **Code** ($K = 5$, range $-1..1$) $\rightarrow$ **Frame**. A continuous underlying attitude is collapsed into a 5-point Likert response; **Frame** measures the mean of the categorical responses, in parallel with a continuous **Sample** $\rightarrow$ **Frame** branch. The loss of information from coarse categorisation becomes visible: ordinal position is recovered, absolute scale is lost — exactly the situation surveyors grapple with.

8.2.5 Cross-tabulation and effect size

Two parallel **Sample** $\rightarrow$ **Code** chains feed **Tab**. Adjusting the correlation between the two underlying **Samples** (by sharing or separating their seeds via **Seed**) moves Cramér’s V smoothly from ≈ 0 (independent) toward 1 (perfectly dependent); the χ^2 test flips from

“fail to reject” to “reject H_0 ” at the same threshold the textbook table predicts. Effect size and inferential decision are visible side-by-side, which addresses a recurring teaching difficulty: students conflate *significant* and *large*.

9 Availability and reproducibility

- **Licence:** GPL-3.0 (free as in freedom).
- **Cost:** free download; VCV Rack itself is also free.
- **Platforms:** macOS (Apple Silicon + Intel), Windows, Linux. No platform-specific code paths; the same C++ compiles for all three via VCV Rack’s official cross-compilation toolchain.
- **Dependencies:** VCV Rack 2 only. All numerical routines (online k-means, Sanger’s PCA, STL decomposition, bootstrap resampling, t-statistic computation via Gaussian approximation of the t-CDF, χ^2 and Cramér’s V computation, autocorrelation estimation) are implemented in the C++ standard library.
- **Installation:** a single `.vcvplugin` archive per plugin, total install size under 1 MB; either dropped into Rack’s user plugins directory, or — once submitted — discoverable through the Rack module library.
- **Reproducibility:** the codebase is short ($\approx 12\,400$ lines C++ across 27 modules in five plugins), reviewable, and ports cleanly between platforms. The same patch loaded on any operating system produces the same behavior at the same parameter settings.

These properties matter for teaching at scale: an instructor can distribute one zip file, send a class of any size to a free download, and trust that every student’s patch will look and behave identically regardless of laptop, OS, or budget.

10 Reproducibility and cross-machine determinism

A central concern for any teaching tool is **reproducibility**: an instructor must be able to share a worked example with students who will see *the same* simulation on their own machines, regardless of operating system or hardware. Empiria addresses this in three layers.

- **Deterministic initialisation.** Every random module in Empiria initializes its RNG from a fixed seed at construction. A fresh patch with default state therefore produces an identical simulation on any machine.
- **Patch-level state preservation.** All module parameters, switch positions, and randomisation choices are saved in the patch JSON. Loading the same patch on another machine reconstructs identical initial conditions.
- **The Seed module.** For experiments in which the user explicitly wishes to *vary* the random realisation (sample a different population, draw a different bootstrap, re-cast a Diffusion S-curve), the dedicated **Seed** module exposes a single integer seed value prominently — knob, RANDOMIZE button, CLOCK-driven increment, a CV output proportional to the value, and a trigger pulse that fires whenever the value changes. The trigger can be patched into any module’s RESET input, so an entire patch can be re-initialized in lockstep. The seed value is part of the patch JSON, so two collaborators who set their Seed knob to the same value see the same simulation.

In **Sample**, common real-world distribution shapes are exposed as right-click presets (adult height, IQ scores, reaction times, survey responses, U-shaped opinion distributions, right-skewed income, etc.), so a student need not tune knobs blindly to set up a recognisable data-generating process.

For *frozen* datasets — situations where the analysis itself should be the variable, not the

data — the **Tape** module records an incoming CV stream into a buffer and replays it deterministically. The same synthetic sample can be passed through Frame, Regress, Test, Boot, and Lag in parallel, or analyzed with different parameter settings, and the underlying data does not change. This is the standard discipline of computational reproducibility, made tactile: a “dataset” lives in the rack as a tangible, scrubbable buffer rather than as an implicit RNG state.

Tape also exposes a right-click “**Export buffer to CSV...**” option that writes the recorded samples to a CSV file (one row per sample, one column per polyphonic channel). The export is immediately loadable in R (`read.csv()`), Python (`pandas.read_csv()`), Julia, or Excel — so a student can carry the synthetic dataset they generated *inside* Empiria into any external analytical environment of their choosing. Combined with the deterministic Seed module, two students running the same patch with the same seed value will export byte-identical CSVs, enabling reproducible assignments in mixed software ecosystems. Patches themselves are small JSON documents (typically 5–50 KB) sharable by email or in a course-management system; built-in VCV Rack screenshot capture (View → Screenshot) produces publication-quality PNGs of the rack at any zoom level.

Pedagogically this matters because it lets *empirical reasoning be shared as a patch file* alongside its data and figures. A small bundle — `.vcv` patch + `.csv` data export + `.png` screenshots — captures the entire context of a worked example in a few hundred kilobytes, sendable to a class of any size, openable on any operating system, and reproducible deterministically.

11 Classroom implementation and proposed evaluation

Empiria is intentionally usable at three different levels of formality, and we sketch deployments for each:

High-school or general-education setting (no prior statistics background). The platform stands on its own as a hands-on demonstration tool — no R, no Python, no command line. A teacher can run a 45-minute session built around a single patch (e.g. `Sample` → `Frame`, sweeping the WIN knob to make the standard error shrink visibly as $1 / \sqrt{n}$), and students engage by turning knobs and predicting what the trace will do. The visual immediacy lowers the prerequisite to roughly the level of a quincunx demonstration, but with vastly more parameter control. Outputs at this level are qualitative: the student says “more samples means less wiggle in the running mean”, and the platform proves it on screen.

Undergraduate introductory statistics or research-methods course. The platform is positioned as a *companion* tool that visualises what R or Python is computing, not as a replacement. Students install VCV Rack (free) and the five Empiria `.vcvplugin` files (total ≈ 1 MB) on personal laptops; no administrator privileges required. A one-class-period orientation introduces the patch-cable interface, the visualization area convention, and four reference patches distributed with the suite (one per pedagogical theme: visible Law-of-Large-Numbers, bootstrap CI, regression with confidence band, agent-based macro–micro composition). Subsequent lecture and lab sessions use Empiria alongside the analytical workflow.

Graduate research-methods or computational-social-science seminar. The same modules become a fast-prototyping sandbox. Students who already know what a t-test *is*

use **Test** to interrogate edge cases (what does the null distribution look like at $df = 3$? what's the actual two-tailed p -value at $t = 1.7$ with $df = 8$ — does the *exact* incomplete-beta calculation disagree with the normal-approximation taught in undergrad?). Polis simulators serve as quick-and-dirty implementations of canonical agent-based models without writing C++; Decisions modules sit beside an actual psychophysics or behavioral-economics dataset, with parameters tuned to fit. CSV export (see *Replicability* below) lets a graduate student carry a synthetic dataset out of the platform and into a methods chapter.

Suggested syllabus integration. A standard three-credit introductory statistics course running 28 sessions can absorb Empiria across approximately ten of those sessions. Sampling distributions and the Central Limit Theorem (sessions 4–5) use **Sample** → **Frame** with the GROWING mode to make the $1/\sqrt{n}$ shrinkage visible. Hypothesis testing (sessions 9–10) uses **Test** with the shaded rejection region as a live demonstration of α and p -values. Confidence intervals and the bootstrap (sessions 12–13) use **Boot** with BCa diagnostics on screen. Simple and multiple regression (sessions 16–18) use **Regress**. Categorical analysis (session 20) uses **Code** → **Tab**. A capstone unit on social-science applications (sessions 25–27) uses Polis, Epi, Space, and Decisions modules to ground statistics in substantive content.

Proposed evaluation design. We propose a within-instructor, between-section randomized comparison over two academic terms. Two parallel sections of the same course, taught by the same instructor, are randomly assigned at the section level (with student consent) to either an Empiria-augmented condition (Empiria plus the usual R workflow) or a control condition (R workflow only, augmented with comparable static figures and Shiny applets to control for the *time-on-visualizations* confound). Both arms receive identical lecture content, identical assessments, and identical homework problems; only the in-class demonstrations differ. Outcomes are measured by (i) the *Comprehensive Assessment of Outcomes*

in a *First Statistics course* (CAOS) administered at the beginning and end of the term, with the *Sampling Variability*, *Statistical Inference*, and *Bivariate Data* subscales as primary endpoints; (ii) a transfer task in which students design (on paper) the inferential pipeline for an unfamiliar applied scenario, scored blind by two graders for completeness and conceptual structure; and (iii) a brief end-of-term survey of self-reported confidence and intrinsic motivation. Mixed-effects models with random intercepts for instructor-section pair and student would handle the nested structure; estimated effect-size targets and minimum sample sizes are reported in the supplementary evaluation pre-registration that accompanies this paper.

Anticipated obstacles. Empiria’s most plausible failure mode is *activation cost*: even though installation is one-click, students who have never opened a modular synthesizer may need 30 minutes of introduction before they can drive the platform productively. The deployment plan above addresses this with a dedicated orientation session and four worked reference patches. A second risk is *content shift*: instructors may be tempted to spend disproportionate time on the visually engaging modules at the expense of analytical content; the suggested syllabus integration above is explicitly *additive* to the usual R workflow rather than substitutive.

12 Architecture

Three design choices deserve note:

- **No external dependencies.** All numerical routines are in the C++ standard library — the plugin builds anywhere VCV Rack builds, with no package-manager overhead. This is what makes the suite pedagogically distributable: a single download per plugin, no install scripts, no Python or R, no conda environments.
- **Audio-thread safety.** Statistics are computed at clock-tick rate, not per-audio-

sample, and cached for the audio thread to read; the UI thread reads the same cache for visualization. Tearing is possible but invisible.

- **Native text rendering.** Module titles and all in-visualization text are drawn through NanoVG in C++ rather than as SVG `<text>` elements, because VCV Rack’s NanoSVG-based panel loader silently drops text from panel SVGs. The shared `ModuleTitle` helper in each plugin’s header centralizes this rendering.

13 Limitations and future work

Test uses the exact Student’s t survival function (regularized incomplete beta) for its p-value, with the null-distribution visualization drawn from the matching t -PDF so the fattening of the tails at low df is visible. **Boot** uses the bias-corrected and accelerated (BCa) interval of Efron (1987) — bias z_0 from the bootstrap distribution, acceleration $\hat{a}$ from the leave-one-out jackknife — both diagnostics surfaced on the panel. **Tab** reports the exact χ^2 upper-tail p-value via the regularized lower incomplete gamma, with significance stars beside it.

The **Network** module publishes its adjacency matrix through VCV Rack’s expander mechanism. Any Empiria simulator placed to its right — currently **Diffusion** (Polis) — automatically switches from well-mixed mass-action to graph-aware transmission and displays $NET \cdot N$ on the panel; the message is also forwarded to any further right neighbour, so additional graph-aware modules can be chained on the same topology. The protocol is documented and any future module (internal or third-party) can opt in by allocating a matching message buffer.

The population cap of 16 agents reflects VCV Rack’s polyphonic-cable maximum and is sufficient for the qualitative dynamics most teaching patches target; lifting it to 32 or 64 would require either splitting state across multiple polyphonic outputs, or rendering the

per-agent state without exposing it on a cable.

Space currently ships three foundational modules; an obvious extension is an Ising/Potts spin model and a Schelling-with-economic- mobility variant. **Decisions** could grow to include a temporal- difference learner, a hyperbolic-discount valuation module, and an ultimatum / dictator-game pair.

The author has used pre-release versions of these modules in his own methodology teaching. Systematic classroom evaluation against a control condition has not yet been carried out and is a natural next step.

14 Acknowledgements

The author thanks the VCV Rack developer community for the modular plugin SDK and example modules, and acknowledges the long open-source tradition of computational social science on which Empiria draws.

15 Disclosure statement

The author reports there are no competing interests to declare. Empiria is released under the GPL-3.0 license through the author’s personal open-source imprint. The imprint is not a commercial entity and yields no income to the author. No funding, university appointment, or commercial sponsorship was tied to the development of the suite; it is a methodological side-project written, documented, and maintained for educational purposes.

16 Data availability statement

This paper describes a software platform rather than the analysis of a specific dataset; the materials that support its results are the source code, panel assets, reference patches, and documentation that together constitute the platform. All such materials are made available under the GPL-3.0 license at the public repository:

[URL WITHHELD FOR DOUBLE-ANONYMOUS REVIEW; the full repository URL is recorded on the title page]

The repository tag `v2.0.0` corresponds to the version of the software described in this paper. The repository contains: (a) the C++17 source for all fifteen Methods modules in `Methods/src/`; (b) panel SVG assets in `Methods/res/`; (c) the three starter patches as portable JSON `.vcv` files under `patches/`; (d) a deterministic patch-generator script `tools/empiria_patch_gen.py` that reproduces the starter patches with any user-supplied seed; (e) the Methods Plugin Manual at `docs/methods_manual.md`; and (f) the source for this paper at `paper.md` together with the bibliography `paper.bib`. The illustrative numerical results reported in the worked examples are themselves outputs of the software at fixed seeds documented in each example, and can be reproduced by any reader who installs the plugin and opens the corresponding `.vcv` file.

References

- American Statistical Association (2016). GAISE college report 2016: Guidelines for assessment and instruction in statistics education.
- Axelrod, R. (1984). *The Evolution of Cooperation*. Basic Books, New York.
- Bass, F. M. (1969). A new product growth for model consumer durables. *Management Science*, 15(5):215–227.

- Belt, A. and VCV Rack contributors (2024). VCV Rack: an open-source modular synthesizer environment.
- Boghosian, B. M. (2014). Kinetics of wealth and the Pareto law. *Physical Review E*, 89(4):042804.
- Bruner, J. S. (1966). *Toward a Theory of Instruction*. Harvard University Press, Cambridge, MA.
- Bush, V. (1931). The differential analyzer. A new machine for solving differential equations. *Journal of the Franklin Institute*, 212(4):447–488.
- Chakraborti, A. (2002). Distributions of money in model markets of economy. *International Journal of Modern Physics C*, 13(10):1315–1321.
- Chance, B., delMas, R., and Garfield, J. (2004). Reasoning about sampling distributions. *The Challenge of Developing Statistical Literacy, Reasoning, and Thinking (Ben-Zvi & Garfield, eds.)*, pages 295–323.
- Cleveland, R. B., Cleveland, W. S., McRae, J. E., and Terpenning, I. (1990). STL: a seasonal-trend decomposition procedure based on Loess. *Journal of Official Statistics*, 6(1):3–73.
- Cobb, G. W. (2007). The introductory statistics course: a Ptolemaic curriculum? *Technology Innovations in Statistics Education*, 1(1).
- Deffuant, G., Neau, D., Amblard, F., and Weisbuch, G. (2000). Mixing beliefs among interacting agents. *Advances in Complex Systems*, 3(01n04):87–98.
- delMas, R., Garfield, J., Ooms, A., and Chance, B. (2007). Assessing students’ conceptual understanding after a first course in statistics. *Statistics Education Research Journal*, 6(2):28–58.

- Efron, B. (1979). Bootstrap methods: another look at the jackknife. *The Annals of Statistics*, 7(1):1–26.
- Efron, B. (1987). Better bootstrap confidence intervals. *Journal of the American Statistical Association*, 82(397):171–185.
- Epstein, J. M. (2006). *Generative Social Science: Studies in Agent-Based Computational Modeling*. Princeton University Press, Princeton, NJ.
- Finzer, W. (2013). The data science education dilemma. *Technology Innovations in Statistics Education*, 7(2).
- Garfield, J. and Ben-Zvi, D. (2009). Helping students develop statistical reasoning: implementing a statistical reasoning learning environment. *Teaching Statistics*, 31(3):72–77.
- Granovetter, M. (1978). Threshold models of collective behavior. *American Journal of Sociology*, 83(6):1420–1443.
- Kermack, W. O. and McKendrick, A. G. (1927). A contribution to the mathematical theory of epidemics. *Proceedings of the Royal Society A*, 115(772):700–721.
- Konold, C. and Lehrer, R. (2008). Technology and mathematics education: an essay in honor of Jim Kaput. *Handbook of International Research in Mathematics Education (English, ed.)*.
- Matsumoto, M. and Nishimura, T. (1998). Mersenne Twister: a 623-dimensionally equidistributed uniform pseudo-random number generator. *ACM Transactions on Modeling and Computer Simulation*, 8(1):3–30.
- Mayer, R. E. (2009). *Multimedia Learning*. Cambridge University Press, Cambridge, 2nd edition.

- Newman, M. E. J. (2002). Spread of epidemic disease on networks. *Physical Review E*, 66(1):016128.
- Press, W. H., Teukolsky, S. A., Vetterling, W. T., and Flannery, B. P. (2007). *Numerical Recipes: The Art of Scientific Computing*. Cambridge University Press, Cambridge, 3rd edition.
- Rogers, E. M. (2003). *Diffusion of Innovations*. Free Press, New York, 5th edition.
- Sanger, T. D. (1989). Optimal unsupervised learning in a single-layer linear feedforward neural network. *Neural Networks*, 2(6):459–473.
- Schelling, T. C. (1971). Dynamic models of segregation. *Journal of Mathematical Sociology*, 1(2):143–186.
- Sfard, A. (1991). On the dual nature of mathematical conceptions: reflections on processes and objects as different sides of the same coin. *Educational Studies in Mathematics*, 22(1):1–36.
- Small, J. S. (2001). *The Analogue Alternative: The Electronic Analogue Computer in Britain and the USA, 1930–1975*. Routledge, London.
- Tintle, N., Chance, B., Cobb, G., Roy, S., Swanson, T., and VanderStoep, J. (2015). Combating anti-statistical thinking using simulation-based methods throughout the undergraduate curriculum. *The American Statistician*, 69(4):362–370.
- Tukey, J. W. (1962). The future of data analysis. *The Annals of Mathematical Statistics*, 33(1):1–67.
- Vergnaud, G. (2009). The theory of conceptual fields. *Human Development*, 52(2):83–94.
- Wilensky, U. (1999). NetLogo. Center for Connected Learning and Computer-Based Modeling, Northwestern University.